SANKET MALI

Full Stack & AI/ML Developer

Pune | +91-9307843884 | sanketsmali27@gmail.com | linkedin.com/in/sanket-mali27 | github.com/SanketMali27

SUMMARY

Final-year B.Tech (CSBS) student building LLM-based systems and full-stack products. Built a LangGraph research agent with RAG, persistent memory, and observability, and deployed a JWT-secured file-sharing app on AWS. Seeking an AI/ML or Full Stack (MERN / Python) internship.

TECHNICAL SKILLS

Languages: Python, JavaScript, C++, SQL

AI / ML: LLMs, RAG, Agentic AI, Multi-Agent Systems, LangChain, LangGraph, MCP, Prompt Engineering, LLM Guardrails, LLM Evaluation, PyTorch, scikit-learn, CNN

Web & Backend: React.js (Vite), Node.js, Express.js, Next.js, FastAPI, REST APIs, JWT, Tailwind CSS, HTML, CSS

Databases: MongoDB, PostgreSQL, SQL, Qdrant (vector DB), Redis

Cloud & Tools: AWS (S3, EC2), Docker, Nginx, Git, GitHub, LangSmith

Data Analysis: Pandas, NumPy, Matplotlib, Jupyter Notebook

PROJECTS

Autonomous Research & Analytics Agent | LangGraph, FastAPI, React, TypeScript, PostgreSQL, Qdrant, Groq, LangSmith

- Built an autonomous AI agent with LangGraph supporting RAG, web search, Python data analysis, and conversational workflows.
- Implemented a RAG pipeline on Qdrant with document ingestion, embeddings, query rewriting, retrieval grading, and answer critique to produce grounded responses.
- Designed persistent chat memory using PostgreSQL, conversation summaries, recent-message context, and LangGraph checkpointing.
- Implemented AI observability and safety using LangSmith, an LLM gateway, guardrails, evaluation, and red-team testing.

CloudShare | React, Node.js, Express, MongoDB, AWS S3, EC2, JWT

- Built a full-stack file-sharing app with JWT access and refresh token authentication and per-user file metadata stored in MongoDB.
- Stored files on AWS S3 with signed URL downloads; added 6-character redeem codes with configurable expiry and download limits, QR-code sharing, and public share pages.
- Deployed on AWS EC2 with PM2 process management and an Nginx reverse proxy, serving the static frontend through Express.

DocuMind AI: Document Classification | Python, PyTorch, Hugging Face, CNN

- Built a PyTorch pipeline that classifies documents as Invoice, Scientific Publication, or Resume.
- Created the data pipeline with Hugging Face Datasets, a custom PyTorch Dataset/DataLoader, image resizing, and tensor conversion.
- Trained a CNN (convolution, ReLU, max-pooling) with CrossEntropyLoss and Adam; used train/validation splits, accuracy evaluation, and best-model checkpointing for inference on unseen documents.

ML Model Comparison System | Python, scikit-learn, Pandas, NumPy, Matplotlib

- Trained and compared 5 classifiers (Logistic Regression, Random Forest, SVM, KNN, Decision Tree) on a Heart Disease dataset under identical preprocessing.
- Built an end-to-end pipeline covering missing-value handling, StandardScaler scaling, train-test split, and automated best-model selection by F1 score.
- Evaluated models on Accuracy, Precision, Recall, and F1; exported results to CSV and visualized them with Matplotlib.

EDUCATION

B.Tech in Computer Science & Business Systems, JSPM Rajarshi Shahu College of Engineering, Pune *2023 – 2027*

CERTIFICATIONS

Full Stack Generative & Agentic AI with Python, Udemy *Feb 2026*

AWS Certified Solutions Architect Associate (Course), Udemy *Mar 2026*

ACHIEVEMENTS

- Solved 400+ problems on LeetCode in C++; participated in hackathons and coding contests.